

HW5

Geoffrey Woollard

My code lives in the repo https://github.com/geoffwoollard/prob_prog

Acknowledgments

For this assignment I primarily acknowledge the patient discussions with Jordan Lovrod, where she held my hand and walked me through the recursive structure of declaring lambda fn's, and the disrinctions between a return from calling something from the `Procedure` class and a `Procedure` object itself. She also supplied code for the classes `Env`, and `Procedure`.

- Ilias Karimalis on the address structure and how it dynamically tracks the evaluations of the dynamic computational graph
- Alan Milligan, Gaurav Bhatt, Masoud Mokhtari and everyone else on Slack for what to do when hitting snags, especially on `empty?`, and `hash-map`, `put` and `get` for handling `str` keys.
- Kim Dinh's comments on the interpretation of problem 1 as a Geometric distribution
- the baseline results provided by Lironne Kurzman
- hw6 starter code for some parts of the `eval_hoppl`, especially the double quotes, and some `primitives`

In [1]:

```
#!/bash
python evaluator.py

FOPPL test 1 passed
FOPPL test 2 passed
FOPPL test 3 passed
FOPPL test 4 passed
FOPPL test 5 passed
FOPPL test 6 passed
FOPPL test 7 passed
FOPPL test 8 passed
FOPPL test 9 passed
FOPPL test 10 passed
FOPPL test 11 passed
FOPPL test 12 passed
FOPPL test 13 passed
HOPPL test 1 passed
HOPPL test 2 passed
HOPPL test 3 passed
HOPPL test 4 passed
HOPPL test 5 passed
HOPPL test 6 passed
HOPPL test 7 passed
HOPPL test 8 passed
HOPPL test 9 passed
HOPPL test 10 passed
HOPPL test 11 passed
HOPPL test 12 passed
All deterministic tests passed
('normal', 5, 1.4142136)
p value 0.9521132603290829
('beta', 2.0, 5.0)
p value 0.8268445434006023
('exponential', 0.0, 5.0)
p value 0.3859297323773355
('normal', 5, 3.2)
p value 0.6247795543285082
('normalmix', 0.1, -1, 0.3, 0.9, 1, 0.3)
p value 0.8588495916821821
('normal', 0, 1.44)
p value 0.6629150112938403
All probabilistic tests passed
1

Sample of prior of program 1:
tensor(56)
2

Sample of prior of program 2:
tensor(-1.5379)
3

Sample of prior of program 3:
tensor([1, 2, 1, 0, 0, 2, 1, 1, 1, 0, 0, 0, 2, 0, 2, 1, 2])
```

Code snippets

In [2]:

```
from dill.source import getsource, getsourcelines
```

In [3]:

```
from evaluator import evaluate, Env, Procedure, standard_env, eval_hoppl

list_of_programs = [evaluate, Env, Procedure, standard_env, eval_hoppl]

for program in list_of_programs:
    for line_number, function_line in enumerate(getsourcelines(program)[0]):
        print(line_number, function_line, end='')
        print()

0 def evaluate(exp, env=None, do_log=False): #TODO: add sigma, or something
1
2     if env is None:
3         env = standard_env()
4
5     fn, _ = eval_hoppl(exp, env, sigma=None, do_log=do_log)
6     ret, sigma = fn(())
7     if do_log: print('return', ret)
8     return ret

0 class Env():
1     "An environment: a dict of {'var': val} pairs, with an outer Env."
2     def __init__(self, params=(), args=(), outer=None):
3         self.data = pmap(zip(params, args))
4         self.outer = outer
5         if outer is None:
6             self.level = 0
7         else:
8             self.level = outer.level+1
9     def __getitem__(self, item):
10         return self.data[item]
11     def find(self, var):
12         "Find the innermost Env where var appears."
13         if (var in self.data):
14             return self
15         else:
16             if self.outer is not None:
17                 return self.outer.find(var)
18             else:
19                 raise RuntimeError('var "{}" not found in outermost scope'.format(var))
20
21     def print_env(self, print_lowest=False):
22         print_line = 1 if print_lowest == False else 0
23         outer = self
24         while outer is not None:
25             if outer.level >= print_limit:
26                 print('Scope on level {}'.format(outer.level))
27                 if 'e' in outer:
28                     print('Found e, ')
29                     print(outer['f'].body)
30                     print(outer['f'].params)
31                     print(outer['f'].env)
32                 print(outer, '\n')
33             outer = outer.outer

0 class Procedure(object):
1     "A user-defined Scheme procedure."
2     def __init__(self, params, body, env, do_log):
3         self.params = self.body, self.env = params, body, env
4         self.do_log = do_log
5     def __call__(self, *args):
6         new_env = copy.deepcopy(self.env)
7         return eval_hoppl(self.body, Env(self.params, args, new_env), do_log=self.do_log)

0 def standard_env():
1     "An Environment with some Scheme standard procedures."
2     env = Env(penv.keys(),, penv.values())
3     return env

0 def eval_hoppl(x, env=standard_env(), sigma=None, do_log=False):
1     # TODO: remove default env=standard_env()
2
3     if do_log: print('x', x)
4
5     if isinstance(x, list):
6         op, param, *args = x
7
8         if op == 'if':
9             assert len(x) == 4
10             test, consequ, alt = x[1:4]
11             if do_log: print('case if: x', x)
12             exp = (conseq if eval_hoppl(test, env, sigma, do_log=do_log)[0] else alt) # be careful to get ret
13             if do_log: print('case if: exp', exp)
14             return eval_hoppl(exp, env, sigma, do_log=do_log)
15
16         if op == 'sample':
17             if do_log: print('case sample: x', x)
18             _, address, exp, distribution = x
19             distribution, sigma = eval_hoppl(exp, distribution, env, sigma, do_log=do_log)
20             if do_log: print('case sample: distribution', distribution)
21             evaluated_sample = distribution.sample()
22             if do_log: print('case sample: evaluated_sample', evaluated_sample)
23             return evaluated_sample, sigma
24
25         elif op == 'observe':
26             if do_log: print('case observe: (pass)')
27             return 'observed', sigma
28         elif op == 'push-address':
29             return '', sigma
30         elif op == 'fn':
31             if do_log: print('case fn: args', args)
32             body = args[0]
33             return Procedure(param, body, env, do_log=do_log), sigma # has eval in it
34             # param ['alpha', 'x']
35             # body [['', ['push-address', 'alpha', 'addr2']], 'x', 'x']
36             # env
37
38         else:
39             if do_log: print('case else: x', x)
40             proc, _ = eval_hoppl(op, env, sigma, do_log=do_log)
41             vals = []
42             if do_log: print('case else: args', args)
43             if do_log: print('case else: args', args)
44             if do_log: print('case else: args', args)
45             if do_log: print('case Procedure: r', r)
46             if do_log: print('case Procedure: r', r)
47             if do_log: print('case primitives: vals[1:]', vals[1:])
48             r = proc(vals[1:]) # primitives
49             if do_log: print('case primitives: r', r)
50
51             return r, sigma
52
53     # base cases
54     elif isinstance(x, str):
55         if x[0] == '"': # daphne output: strings have double, double quotes
56             return x[1:-1], sigma
57         lowest_env = env.find(x)
58         return lowest_env(x), sigma
59
60     elif isinstance(x, (float, int, bool)):
61         return torch.tensor(x), sigma
62
63     else:
64         raise ValueError('unknown expression case')
```

In [4]:

```
import primitives

for line_number, function_line in enumerate(getsourcelines(primitives)[0]):
    print(line_number, function_line, end='')

0 """primitives (and distributions) used to evaluate algorithm 6 in
1 van de Meent, J.-W., Paige, B., Yang, H., & Wood, F. (2018).
2 An Introduction to Probabilistic Programming, XX(Xx), 1-221.
3 http://doi.org/10.1561/XXXXXXX
4
5 Acknowledgements to Yuan T https://github.com/yuant95/CPSC532W/blob/master/HW2/primitives.py
6 and
7 Masoud Mokhtari https://github.com/MasoudMo/CPSC-532W/blob/master/HW2/primitives.py
8 and
9 RM6 starter code https://www.cs.ubc.ca/~fwood/CS532W-539W/homework/6.html
10 """
11
12 import torch
13 from torch import tensor
14
15 number = (float, int)
16 distribution_types = (
17     torch.distributions.Normal,
18     torch.distributions.Beta,
19     torch.distributions.Uniform,
20     torch.distributions.Exponential,
21     torch.distributions.Categorical,
22     torch.distributions.bernoulli.Bernoulli,
23     torch.distributions.dirichlet.Dirichlet,
24     torch.distributions.gamma.Gamma,
25 )
26
27
28 def two_arg_op_primitive(op, arg1, arg2):
29     arg1, arg2 = arg1, arg2
30     return op(arg1, arg2)
31
32
33 def add_primitive(arg1, arg2):
34     return two_arg_op_primitive(torch.add, arg1, arg2)
35
36
37 def subtract_primitive(arg1, arg2):
38     return two_arg_op_primitive(torch.subtract, arg1, arg2)
39
40
41 def multiply_primitive(arg1, arg2):
42     return two_arg_op_primitive(torch.multiply, arg1, arg2)
43
44
45 def divide_primitive(arg1, arg2):
46     return two_arg_op_primitive(torch.divide, arg1, arg2)
47
48
49 def one_arg_op_primitive(op, arg):
50     arg0 = arg[0] # because list of len one passed, i.e. [arg0]
51     return op(arg0)
52
53
54 def sqrt_primitive(arg):
55     return one_arg_op_primitive(torch.sqrt, arg)
56
57
58 def get_primitive(vector_and_index):
59     vector, index = vector_and_index
60
61     if isinstance(index, str):
62         index_safe = index
63     elif torch.is_tensor(index):
64         index_safe = index.item()
65     else:
66         assert False, 'index type {} case not implemented'.format(type(index))
67
68     if isinstance(vector, dict):
69         return vector[index_safe]
70     elif torch.is_tensor(vector):
71         assert not isinstance(index, str)
72         return vector[index.long()]
73     elif isinstance(vector, list):
74         index_int = torch.tensor(int(index))
75         assert torch.isclose(index_int, index) # TODO: use native pytorch
76         return vector[index_int]
77     else:
78         assert False, 'vector type {} case not implemented'.format(type(vector))
79
80
81 def put_primitive(vector, index, overwritevalue):
82     vector, index, overwritevalue = vector, index, overwritevalue
83
84     if isinstance(index, str):
85         index_safe = index
86     elif torch.is_tensor(index):
87         index_safe = index.item()
88     else:
89         assert False, 'index type {} case not implemented'.format(type(index))
90
91     if isinstance(vector, dict):
92         vector[index_safe] = overwritevalue
93     elif torch.is_tensor(vector):
94         assert not isinstance(index, str)
95         vector[index.long()] = overwritevalue
96     else:
97         assert False, 'vector type {} case not implemented'.format(type(vector))
98     return vector
99
100
101 def return_idx_primitive(vector, idx, idx_f):
102     return vector[0][idx, idx_f]
103
104
105 def first_primitive(vector):
106     return return_idx_primitive(vector, idx=0, idx_f=1)
107
108
109 def second_primitive(vector):
110     return return_idx_primitive(vector, idx=1, idx_f=2)
111
112
113 def last_primitive(vector):
114     return return_idx_primitive(vector, idx=-1, idx_f=None)
115
116
117 def nth_primitive(vector, nth):
118     vector, nth = vector, nth
119     return return_idx_primitive(vector, idx=-nth, idx_f=-nth+1)
120
121
122 def hash_map_primitive(hash_pairs):
123     keys = hash_pairs[:, :2]
124     for idx, key in enumerate(keys):
125         if torch.is_tensor(key):
126             tensor_key = key
127             keys[idx] = tensor_key.item() # dict keys as tensors problematic. can make but lookup fails
128             # dict keys as tensors problematic. can make but lookup fails
129             if isinstance(key, str):
130                 keys[idx] = key # if key string, just keep as is
131
132     # keys = [tensor_key.item() for tensor_key in keys]
133     vals = hash_pairs[:, 2:]
134     return dict(zip(keys, vals))
135
136
137 def gt_primitive(consequent_alternative):
138     return two_arg_op_primitive(torch.gt, consequent_alternative)
139
140
141 def lt_primitive(consequent_alternative):
142     return two_arg_op_primitive(torch.lt, consequent_alternative)
143
144
145 def ge_primitive(consequent_alternative):
146     return two_arg_op_primitive(torch.ge, consequent_alternative)
147
148
149 def le_primitive(consequent_alternative):
150     return two_arg_op_primitive(torch.le, consequent_alternative)
151
152
153 def eq_primitive(consequent_alternative):
154     return two_arg_op_primitive(torch.eq, consequent_alternative)
155
156
157 def rest_primitive(vector):
158     return vector[0][1:]
159
160
161 def freshvar_primitive(arg):
162     return None
163
164
165 def vector_primitive(vector):
166     ret = list()
167     for e in vector:
168         try:
169             ret.append(e.tolist())
170         except:
171             ret.append(e)
172     try:
173         return torch.tensor(ret)
174     except:
175         return ret
176
177
178 def append_primitive(vector, element):
179     vector, element = vector, element
180     return torch.cat((vector, torch.Tensor([element])), 0)
181
182
183 def tanh_primitive(arg):
184     return one_arg_op_primitive(torch.tanh, arg)
185
186
187 def and_primitive(arg1, arg2):
188     return two_arg_op_primitive(torch.logical_and, arg1, arg2)
189
190
191 def or_primitive(arg1, arg2):
192     return two_arg_op_primitive(torch.logical_or, arg1, arg2)
193
194
195 def abs_primitive(arg):
196     return one_arg_op_primitive(torch.abs, arg)
197
198
199 def empty_primitive(args):
200     vector = args[0]
201     if torch.is_tensor(vector) or isinstance(vector, list):
202         return len(vector) == 0
203     else:
204         assert False, 'length for non list or non tensor not implemented'
205
206
207 def cons_primitive(args):
208     """https://bfontaine.net/blog/2014/05/25/how-to-remember-the-difference-between-conj-and-cons-in-clojur
209     """
210     item, vector = args
211     if torch.is_tensor(item) and torch.is_tensor(vector):
212         return torch.cat((torch.tensor(item), vector), dim=0)
213     elif isinstance(vector, list):
214         return [item] + vector
215     else:
216         assert False, 'not implemented'
217
218
219 def prepend_primitive(args):
220     """https://bfontaine.net/blog/2014/05/25/how-to-remember-the-difference-between-conj-and-cons-in-clojur
221     """
222     vector, item = args
223     if torch.is_tensor(item) and torch.is_tensor(vector):
224         if item.dim() == 0:
225             item = item.reshape(1,)
226         elif item.dim() == 1:
227             pass
228         else:
229             assert False, 'not implemented'
230         return torch.cat((item, vector), dim=0)
231     elif isinstance(vector, list):
232         return vector + [item]
233     else:
234         assert False, 'not implemented'
235
236
237 def conj_primitive(args):
238     """https://bfontaine.net/blog/2014/05/25/how-to-remember-the-difference-between-conj-and-cons-in-clojur
239     """
240     vector, item = args
241     if torch.is_tensor(item) and torch.is_tensor(vector):
242         assert item.dim() == 0
243         return torch.cat((vector, item.reshape(1,)), dim=0)
244     elif isinstance(vector, list):
245         return vector + [item]
246     else:
247         assert False, 'not implemented'
248
249
250 def log_primitive(arg):
251     return one_arg_op_primitive(torch.log, arg)
252
253
254 def peek_primitive(vector):
255     vector = vector[0]
256     # TODO: assert only defined for vectors
257     return vector[0]
258
259
260 # NB: these functions take a list [c0] or [c0, c1, ..., cn]
261 # rely on user to not write something non-sobal that will fail (e.g. ["+", 1, 2, 3])
262 primitives_d = {
263     '+': add_primitive,
264     '-': subtract_primitive,
265     '/': divide_primitive,
266     '*': multiply_primitive,
267     'sqrt': sqrt_primitive,
268     'vector': vector_primitive,
269     'get': get_primitive,
270     'put': put_primitive,
271     'first': first_primitive,
272     'second': second_primitive,
273     'last': last_primitive,
274     'nth': nth_primitive,
275     'append': append_primitive,
276     'hash-map': hash_map_primitive,
277     '>': gt_primitive,
278     '<': lt_primitive,
279     '>=': ge_primitive,
280     '<=': le_primitive,
281     'rest': rest_primitive,
282     'mat-transpose': lambda a: a[0].T,
283     'mat-mul': lambda a: torch.matmul(a[0], a[1]),
284     'mat-add': add_primitive,
285     'mat-repeat': lambda a: a[0].repeat((int(a[1].item()), int(a[2].item()))),
286     'and': and_primitive,
287     'or': or_primitive,
288     'abs': abs_primitive,
289     'empty?': empty_primitive,
290     'cons': cons_primitive,
291     'conj': conj_primitive,
292     'log': log_primitive,
293     'peek': peek_primitive,
294     'prepend': prepend_primitive,
295 }
296
297
298 def normal(mean_std):
299     return two_arg_op_primitive(torch.distributions.Normal, mean_std)
300
301
302 def beta(alpha_beta):
303     return two_arg_op_primitive(torch.distributions.Beta, alpha_beta)
304
305
306 def exponential(lam):
307     return one_arg_op_primitive(torch.distributions.Exponential, lam)
308
309
310 def uniform(low_hi):
311     return two_arg_op_primitive(torch.distributions.Uniform, low_hi)
312
313
314 def discrete(prob_vector):
315     return one_arg_op_primitive(torch.distributions.Categorical, prob_vector)
316
317
318 def flip(prob):
319     return one_arg_op_primitive(torch.distributions.bernoulli.Bernoulli, prob)
320
321
322 def dirichlet(concentration):
323     return one_arg_op_primitive(torch.distributions.dirichlet.Dirichlet, concentration)
324
325
326 def gamma(concentration_rate):
327     return two_arg_op_primitive(torch.distributions.gamma.Gamma, concentration_rate)
328
329
330 distributions_d = {
331     'normal': normal,
332     'beta': beta,
333     'exponential': exponential,
334     'uniform-continuous': uniform,
335     'discrete': discrete,
336     'flip': flip,
337     'dirichlet': dirichlet,
338     'gamma': gamma,
339 }
340
341 penv = (**distributions_d, **primitives_d)
```

Program 2 (1.daphne)

```
In [54]: from evaluator import evaluate, ast_helper
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
In [48]: i=1
fname='{}.daphne'.format(i)
exp = ast_helper(fname,directory='programs/')
%cat programs/1.daphne
```

```
(defn until-success [p n]
  (if (sample (flip p))
      n
      (until-success p (+ n 1))))

(let [p 0.01]
  (until-success p 0))
```

```
In [85]: evaluate(exp, do_log=False) # example return value
```

Out[85]: tensor(18)


```
In [45]: import sys
sys.setrecursionlimit(1000000)
```

```
In [ ]: n_samples=100000
samples = [evaluate(exp).item() for sample in range(n_samples)]
# 4.8s / 100 samples
```

```
In [84]: # np.save('program2.npy',np.array(samples))
```

```
In [72]: sr = pd.Series(samples)
sr.plot.hist(bins=50)
plt.xlabel('flips until success')
plt.title('{} \n n_samples {}'.format(fname,n_samples))
```

Out[72]: Text(0.5, 1.0, '2.daphne \n n_samples 100000')



```
In [73]: print('expectation w.r.t. the prior {:.13f}'.format(sr.mean()))
print('std & var w.r.t. the prior {:.13f} & {:.1f}'.format(sr.std(),sr.var()))
```

expectation w.r.t. the prior 98.665
std & var w.r.t. the prior 99.041 & 9809.1

This is a standard textbook problem of a [Geometric distribution](#) "The probability distribution of the number $Y = X - 1$ of failures before the first success, supported on the set $\{0, 1, 2, \dots\}$."

The ground truth mean and var are thus $\frac{1-p}{p}$ and $\frac{1-p}{p^2}$, where $p = 0.01$ in the homework problem, and we can analytically compare against our estimates.

```
In [83]: p = 0.01
gt_mean = (1-p)/p
gt_std = np.sqrt(gt_mean/p)
gt_mean, gt_std
assert np.abs(gt_mean - sr.mean()) / gt_mean < 0.05
assert np.abs(gt_std - sr.std()) / gt_std < 0.05
```


Program 3 (2.daphne)

```
In [17]: from evaluator import evaluate, ast_helper
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
In [3]: i=2
fname='{}.daphne'.format(i)
exp = ast_helper(fname,directory='programs/')
%cat programs/2.daphne

(defn marsaglia-normal [mean var]
  (let [d (uniform-continuous -1.0 1.0)
        x (sample d)
        y (sample d)
        s (+ (* x x) (* y y))]
    (if (< s 1)
      (+ mean (* (sqrt var)
                  (* x (sqrt (* -2 (/ (log s) s))))))
      (marsaglia-normal mean var))))

(let [mu (marsaglia-normal 1 5)
      sigma (sqrt 2)
      lik (normal mu sigma)]
  (observe lik 8)
  (observe lik 9)
  mu)
```

```
In [21]: evaluate(exp, do_log=False) # example of the return
```

Out[21]: tensor(2.9357)


```
In [6]: import sys
sys.setrecursionlimit(1000000)
```

```
In [ ]: n_samples=1000*200
samples = [evaluate(exp).item() for sample in range(n_samples)]
# 10.2s / 1000 samples to 200k in 30 min
```

```
In [22]: # np.save('program3.npy',np.array(samples))
```

```
In [15]: sr = pd.Series(samples)
sr.plot.hist(bins=50)
plt.xlabel('mu')
plt.title('{} \n n_samples {}'.format(fname,n_samples))
```

Out[15]: Text(0.5, 1.0, '3.daphne \n n_samples 200000')



```
In [14]: print('expectation w.r.t. the prior {:.3f}'.format(sr.mean()))
print('std & var w.r.t. the prior {:.3f} & {:.1f}'.format(sr.std(),sr.var()))
```

expectation w.r.t. the prior 1.003
std & var w.r.t. the prior 2.239 & 5.0

The program follows its namesake, the [Marsaglia polar method](#), and so we know the distribution of the prior is $\mathcal{N}[\mu|0, 5]$. We can thus check that we are within 5% tolerance.

```
In [20]: gt_mean = 1
gt_var = 5

assert np.abs(gt_mean - sr.mean()) / gt_mean < 0.05
assert np.abs(gt_var - sr.var()) / gt_var < 0.05
```

A normally distributed random quantity, via transformation and rejection. Take a little time to think about the sampled values of x and y and be amazed that this works. Think a little about how to deal with this kind of case in amortized inference settings.

The prior is using control flow with the `if` statement, essentially rejection sampling to ensure (x,y) are "inside the unit circle".

So we can get a new type of distribution, a normal with two parameters, from just two continuous distributions x and y.

In amortized inference, we could use some NN transformation of uniform RVs, to learn a new distribution. If the distribution was normal, perhaps they would learn the Marsaglia polar method, or some other method to sample from a Normal (the wiki page mentioned a few).

If our posterior was some arbitrary distribution, we could fit the parameters of the NN in amortized inference, and learn the custom transform for that posterior.

Program 4 (3.daphne)

```
In [24]: from evaluator import evaluate, ast_helper
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
In [3]: i=3
fname='{}.daphne'.format(i)
exp = ast_helper(fname,directory='programs/')
%cat programs/3.daphne

(defn reduce [f x values]
  (if (empty? values)
      x
      (reduce f (f x (first values)) (rest values))))

(let [observations [0.9 0.8 0.7 0.0 -0.025 -5.0 -2.0 -0.1 0.0 0.13 0.45 6 0.2 0.3 -1 -1]
      init-dist (discrete [1.0 1.0 1.0])
      trans-dists {0 (discrete [0.1 0.5 0.4])
                   1 (discrete [0.2 0.2 0.6])
                   2 (discrete [0.15 0.15 0.7])}
      obs-dists {0 (normal -1 1)
                 1 (normal 1 1)
                 2 (normal 0 1)})

  (reduce
   (fn [states obs]
     (let [state (sample (get trans-dists
                           (peek states)))]
       (observe (get obs-dists state) obs)
       (conj states state)))
   [(sample init-dist)]
   observations))
```

It's an HMM again, this time implemented generically. Take the time to read this source code to see how this works. When you get this working you can be very proud. You will be most of the way towards a very powerful HOPPL language implementation.

This program works by making use of a reduce over a function that does the HMM step. I.e.

- f in defn reduce [f x values] is (fn [states obs] ... (conj states state)))
- x is [(sample init-dist)]
- values is observations

Furthermore we can include a read for the observatoins, and don't have to inline that.

The reduce module at the end is also modular to any sized problem, not just the 3 states here. We could have init-dist, trans-dists, and obs-dists on disk and read them in, and the reduce module would still work on them.

```
In [4]: evaluate(exp, do_log=False) # example of the return
```

```
Out[4]: tensor([2, 2, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 1, 2, 2, 2, 2])
```

```
In [6]: import sys
sys.setrecursionlimit(1000000)
```


```
In [ ]: n_samples=100*190
samples = [evaluate(exp).tolist() for sample in range(n_samples)]
# 9.4s / 100 samples
```

```
In [14]: samples_array = np.array([sample.tolist() for sample in samples])
# np.save('program4.npy',samples_array)
```

```
In [29]: df = pd.DataFrame(samples_array)
df_wide = pd.melt(df.reset_index(),id_vars='index')
```

```
In [37]: plt.figure(figsize=(15,5))
ax = sns.countplot(x="value", hue="variable", data=df_wide)
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0)
```

```
Out[37]: <matplotlib.legend.Legend at 0x1396f5490>
```



```
In [47]: mean = samples_array.mean(0)
std = samples_array.std(0)
var = samples_array.var(0)

for idx in range(samples_array.shape[1]):
    print('dim {} w.r.t. the prior | expectation {:.13f} | std {:.13f} | var {:.13f}'.format(idx, mean[idx], var
```

```
dim 0 w.r.t. the prior | expectation 0.991 | std 0.667 | var 0.817
dim 1 w.r.t. the prior | expectation 1.419 | std 0.542 | var 0.736
dim 2 w.r.t. the prior | expectation 1.415 | std 0.539 | var 0.734
dim 3 w.r.t. the prior | expectation 1.415 | std 0.542 | var 0.736
dim 4 w.r.t. the prior | expectation 1.420 | std 0.535 | var 0.732
dim 5 w.r.t. the prior | expectation 1.416 | std 0.538 | var 0.734
dim 6 w.r.t. the prior | expectation 1.416 | std 0.539 | var 0.734
dim 7 w.r.t. the prior | expectation 1.427 | std 0.536 | var 0.732
dim 8 w.r.t. the prior | expectation 1.423 | std 0.542 | var 0.736
dim 9 w.r.t. the prior | expectation 1.404 | std 0.549 | var 0.741
dim 10 w.r.t. the prior | expectation 1.420 | std 0.540 | var 0.735
dim 11 w.r.t. the prior | expectation 1.416 | std 0.546 | var 0.739
dim 12 w.r.t. the prior | expectation 1.414 | std 0.540 | var 0.735
dim 13 w.r.t. the prior | expectation 1.426 | std 0.535 | var 0.732
dim 14 w.r.t. the prior | expectation 1.413 | std 0.546 | var 0.739
dim 15 w.r.t. the prior | expectation 1.421 | std 0.535 | var 0.731
dim 16 w.r.t. the prior | expectation 1.410 | std 0.545 | var 0.738
```