

HW6

Geoffrey Woollard

My code lives in the repo https://github.com/geoffwoollard/prob_prog

Acknowledgments

- discussions with Jordan Lovrod, Ilias Karimalis, Gaurav Bhatt
- [starter code](#) for `smc.resample_particles` and `smc.SMC` from Masoud Mokhtari

```
In [1]: from dill.source import getsource, getsourcelines

In [2]: from smc import resample_particles, SMC

list_of_programs = [resample_particles, SMC]

for program in list_of_programs:
    for line_number, function_line in enumerate(getsourcelines(program)[0]):
        print(line_number, function_line, end='')
    print()

0 def resample_particles(particles, log_weights):
1     """
2     Eq. 4.24 in course textbook (https://arxiv.org/abs/1809.10756v2, pp. 122)
3     See Algorithm 15 in the course textbook, section 6.7 Sequential Monte Carlo, p. 176
4     """
5     log_weights = tensor(log_weights)
6     n_particles = log_weights.size().numel()
7
8     unnormalized_particle_weights = torch.exp(log_weights).detach().numpy()
9
10    particle_idx = np.random.choice(
11        a=range(n_particles),
12        size=n_particles,
13        p=unnormalized_particle_weights/unnormalized_particle_weights.sum(),
14        replace=True,
15    )
16    #print('particle_idx',particle_idx)
17
18    new_particles = []
19    for idx in range(n_particles):
20        new_particles.append(particles[particle_idx[idx]]) # TODO: copy?
21
22    log_Z = np.log(np.sum(unnormalized_particle_weights)/n_particles)
23    return log_Z, new_particles

0 def SMC(n_particles, exp,do_log=False):
1
2     particles = []
3     weights = []
4     logZs = []
5     output = lambda x: x
6
7     for i in range(n_particles):
8
9         res = evaluate(exp, env=None)('addr_start', output)
10        logW = 0.
11
12
13        particles.append(res)
14        weights.append(logW)
15
16        #can't be done after the first step, under the address transform, so this should be fine:
17        done = False
18        smc_cnter = 0
19        while not done:
20            if do_log: print('In SMC step {}, Zs: '.format(smc_cnter), logZs)
21            for i in range(n_particles): #Even though this can be parallelized, we run it serially
22                res = run_until_observe_or_end(particles[i]) # particle i at next breakbpoint
23                if 'done' in res[2]: #this checks if the calculation is done
24                    particles[i] = res[0]
25                    if i == 0:
26                        done = True #and enforces everything to be the same as the first particle
27                        address = ''
28                else:
29                    if not done: # triggered when i=0 was not done and i>0 was done
30                        raise RuntimeError('Failed SMC, finished one calculation before the other')
31            else:
32                #TODO: check particle addresses, and get weights and continuations
33                particles[i] = res
34                cont, args, sigma = res
35                assert 'observe' == sigma['type']
36                weights[i] = sigma['distribution'].log_prob(sigma['observed_constant'])
37
38                # check particle addresses
39                if i == 0:
40                    break_point_address = sigma['address']
41                else:
42                    if sigma['address'] != break_point_address:
43                        assert False, 'particles at different break points'
44
45
46
47            if not done:
48                #resample and keep track of logZs
49                logZn, particles = resample_particles(particles, weights)
50                logZs.append(logZn)
51            smc_cnter += 1
52    logZ = sum(logZs)
53    return logZ, particles
```